

Practical 1

Aim: Implementation of logic programming using Prolog (DFS) – Water Jug Problem

Code: water_jug.pl

```
water_jug(X, Y) :-  
    X > 4, Y < 3, write('4L jug overflow. '), nl.  
  
water_jug(X, Y) :-  
    X < 4, Y > 3, write('3L jug overflow. '), nl.  
  
water_jug(X, Y) :-  
    X > 4, Y > 3, write('Both jugs overflow. '), nl.  
  
water_jug(X, Y) :-  
    (X == 0, Y == 0, nl, write('4L:0 & 3L:3 (Action: Fill 3L jug.)'), YY is 3,  
    water_jug(X, YY));  
  
    (X == 0, Y == 0, nl, write('4L:4 & 3L:0 (Action: Fill 4L jug.)'), XX is 4,  
    water_jug(XX, Y));  
  
    (X == 2, Y == 0, nl, write('4L:2 & 3L:0 (Action: Goal State reached...)' ));  
  
    (X == 4, Y == 0, nl, write('4L:1 & 3L:3 (Action: Pour water from 4L to 3L jug.)'),  
    XX is X - 3, YY is 3, water_jug(XX, YY));  
  
    (X == 0, Y == 3, nl, write('4L:3 & 3L:0 (Action: Pour water from 3L to 4L jug.)'),  
    XX is 3, YY is 0, water_jug(XX, YY));  
  
    (X == 1, Y == 3, nl, write('4L:1 & 3L:0 (Action: Empty 3L jug.)'), YY is 0,  
    water_jug(X, YY));  
  
    (X == 3, Y == 0, nl, write('4L:3 & 3L:3 (Action: Fill 3L jug.)'), YY is 3,  
    water_jug(X, YY));  
  
    (X == 3, Y == 3, nl, write('4L:4 & 3L:2 (Action: Pour water from 3L jug to 4L jug until 4L jug  
is full.)'),  
    XX is X + 1, YY is Y - 1, water_jug(XX, YY));  
  
    (X == 1, Y == 0, nl, write('4L:0 & 3L:1 (Action: Pour water from 4L jug to 3L jug.)'),  
    XX is Y, YY is X, water_jug(XX, YY));  
  
    (X == 0, Y == 1, nl, write('4L:4 & 3L:1 (Action: Fill 4L jug.)'),  
    XX is 4, water_jug(XX, Y));
```

Output:

```
SWI-Prolog (AMD64, Multi-threaded, version 9.2.9)
File Edit Settings Run Debug Help
Welcome to SWI-Prolog (threaded, 64 bits, version 9.2.9)
SWI-Prolog comes with ABSOLUTELY NO WARRANTY. This is free software.
Please run ?- license. for legal details.

For online help and background, visit https://www.swi-prolog.org
For built-in help, use ?- help(Topic). or ?- apropos(Word).

?-
% d:/mohd zaidd shaikh (124) aial/water_jug compiled 0.00 sec. 0 clauses
?- water_jug(4, 3)
|
false.

?- water_jug(2, 3).
4L:2 & 3L:0 (Action: Empty 3L jug.)
4L:2 & 3L:0 (Action: Goal State reached...)
true ■
```

Practical 2

Aim: Implementation of logic programming using Prolog (BFS) – Tic Tac Toe

Code: tic-tac-toe.pl

```
% Minimal Tic Tac Toe game in Prolog (2-player, terminal-based)

% Initial empty board
board([' ', ' ', ' ',
      ' ', ' ', ' ',
      ' ', ' ', ' ']).

% Display the board
display_board([A,B,C,D,E,F,G,H,I]) :-
    format('~w | ~w | ~w~n', [A,B,C]),
    format('--+---+~n'),
    format('~w | ~w | ~w~n', [D,E,F]),
    format('--+---+~n'),
    format('~w | ~w | ~w~n~n', [G,H,I]).

% Make a move: replace N-th position (1-indexed) with X or O
move(Board, Pos, Player, NewBoard) :-
    nth1(Pos, Board, ' '),      % Ensure the spot is empty
    replace(Board, Pos, Player, NewBoard).

% Replace helper
replace([_|T], 1, X, [X|T]).
replace([H|T], I, X, [H|R]) :-
    I > 1, I1 is I - 1, replace(T, I1, X, R).

% Win conditions
win(Board, Player) :-
    member([A,B,C], [[1,2,3], [4,5,6], [7,8,9],
                     [1,4,7], [2,5,8], [3,6,9],
                     [1,5,9], [3,5,7]]),
    nth1(A, Board, Player),
    nth1(B, Board, Player),
    nth1(C, Board, Player).

% Start game
play :-
    board(B), display_board(B),
    play_turn(B, 'X').

% Alternate turns
play_turn(Board, Player) :-
```

```
write(Player), write("'s turn. Enter position (1-9): "),
read(Pos),
move(Board, Pos, Player, NewBoard),
display_board(NewBoard),
( win(NewBoard, Player) ->
    write(Player), write(' wins!'), nl
; switch(Player, Next), play_turn(NewBoard, Next)
).

% Switch player
switch('X', 'O').
switch('O', 'X').
```

Output:

```

SWI-Prolog (AMD64, Multi-threaded, version 9.2.9)
File Edit Settings Run Debug Help
SWI-Prolog comes with ABSOLUTELY NO WARRANTY. This is free software.
Please run ?- license. for legal details.

For online help and background, visit https://www.swi-prolog.org
For built-in help, use ?- help(Topic). or ?- apropos(Word).

?-
% d:/mohd zaidd shaikh (124) aial/tic-tac-toe compiled 0.00 sec, -3 clauses
?- play.
  |  | 
--+--+--+
  |  | 
--+--+--+
  |  | 

X's turn. Enter position (1-9): 1
|: .
X |  | 
--+--+--+
  |  | 
--+--+--+
  |  | 

O's turn. Enter position (1-9): |: 3.
X |  | O
--+--+--+
  |  | 
--+--+--+
  |  | 

X's turn. Enter position (1-9): |: 2.
X | X | O
--+--+--+
  |  | 
--+--+--+
  |  | 

O's turn. Enter position (1-9): |: 5.
X | X | O
--+--+--+
  | O | 
--+--+--+
  |  | 

X's turn. Enter position (1-9): |: 4
|: .
X | X | O
--+--+--+
X | O | 
--+--+--+
  |  | 

O's turn. Enter position (1-9): |: 7.
X | X | O
--+--+--+
X | O | 
--+--+--+
O |  | 

O wins!
true .
?- █

```

Practical 3

Aim: Implementation of logic programming using Prolog (Hill Climbing) – 8 Puzzle

Code: 8_puzzle.pl

```
% Simple Prolog Planner for the 8 Puzzle Problem

/* This predicate initialises the problem states. The first argument of solve is the initial state,
the 2nd the goal state, and the third the plan that will be produced.*/

test(Plan):-
    write('Initial state:'),nl,
    Init= [at(tile4,1), at(tile3,2), at(tile8,3), at(empty,4), at(tile2,5), at(tile6,6), at(tile5,7),
at(tile1,8), at(tile7,9)],
    write_sol(Init),
    Goal= [at(tile1,1), at(tile2,2), at(tile3,3), at(tile4,4), at(empty,5), at(tile5,6), at(tile6,7),
at(tile7,8), at(tile8,9)],
    nl,write('Goal state:'),nl,
    write(Goal),nl,nl,
    solve(Init,Goal,Plan).

solve(State, Goal, Plan):-
    solve(State, Goal, [], Plan).

% Determines whether Current and Destination tiles are a valid move.
is_movable(X1,Y1) :- (1 is X1 - Y1) ; (-1 is X1 - Y1) ; (3 is X1 - Y1) ; (-3 is X1 - Y1).

/* This predicate produces the plan. Once the Goal list is a subset of the current State the
plan is complete and it is written to the screen using write_sol */

solve(State, Goal, Plan, Plan):-
    is_subset(Goal, State), nl,
    write_sol(Plan).

solve(State, Goal, Sofar, Plan):-
    act(Action, Preconditions, Delete, Add),
    is_subset(Preconditions, State),
    \+ member(Action, Sofar),
    delete_list(Delete, State, Remainder),
    append(Add, Remainder, NewState),
    solve(NewState, Goal, [Action|Sofar], Plan).

/* The problem has three operators.
1st arg = name
2nd arg = preconditions
```

```
3rd arg = delete list
4th arg = add list. */

% Tile can move to new position only if the destination tile is empty & Manhattan distance = 1
act(move(X,Y,Z),
    [at(X,Y), at(empty,Z), is_movable(Y,Z)],
    [at(X,Y), at(empty,Z)],
    [at(X,Z), at(empty,Y)]).

% Utility predicates.
% Check is first list is a subset of the second
is_subset([H|T], Set):-
    member(H, Set),
    is_subset(T, Set).
is_subset([], _).
% Remove all elements of 1st list from second to create third.
delete_list([H|T], Curstate, Newstate):-
    remove(H, Curstate, Remainder),
    delete_list(T, Remainder, Newstate).
delete_list([], Curstate, Curstate).

remove(X, [X|T], T).
remove(X, [H|T], [H|R]):-
    remove(X, T, R).

write_sol([]).
write_sol([H|T]):-
    write_sol(T),
    write(H), nl.

append([H|T], L1, [H|L2]):-
    append(T, L1, L2).
append([], L, L).

member(X, [X|_]).
member(X, [_|T]):-
    member(X, T).
```

Output:

```
SWI-Prolog (AMD64, Multi-threaded, version 9.2.9)
File Edit Settings Run Debug Help
Welcome to SWI-Prolog (threaded, 64 bits, version 9.2.9)
SWI-Prolog comes with ABSOLUTELY NO WARRANTY. This is free software.
Please run ?- license. for legal details.

For online help and background, visit https://www.swi-prolog.org
For built-in help, use ?- help(Topic). or ?- apropos(Word).

?-
% d:\mohd zaid shaikh (124) aial\8_puzzle compiled 0.00 sec. -3 clauses
?- test(Plan).
Initial state:
at(tile7,9)
at(tile1,8)
at(tile5,7)
at(tile6,6)
at(tile2,5)
at(empty,4)
at(tile8,3)
at(tile3,2)
at(tile4,1)

Goal state:
[at(tile1,1).at(tile2,2).at(tile3,3).at(tile4,4).at(empty,5).at(tile5,6).at(tile6,7).at(tile7,8).at(tile8,9)]

false.
?- █
```


Practical 4

Aim: Intro to Python Libraries – Basic Libraries, NumPy, Pandas

#numpy

```
[3] x=np.array([2,3,4,5])  
    print(type(x))
```

```
... <class 'numpy.ndarray'>
```

```
[5] print(x)
```

```
... [2 3 4 5]
```

```
[7] x=np.array([2,3,'n',5])  
    print(x)
```

```
... ['2' '3' 'n' '5']
```

```
[9] d=np.arange(start=1, stop=10, step=2)  
    print(d)
```

```
... [1 3 5 7 9]
```

```
grid = np.arange(start =1, stop=10).reshape(3,3)
print(grid)
```

[23]

```
...  [[1 2 3]
      [4 5 6]
      [7 8 9]]
```

```
np.ones((3,4))
```

[11]

```
...  array([[1., 1., 1., 1.],
          [1., 1., 1., 1.],
          [1., 1., 1., 1.]])
```

```
np.random.rand(5)
```

[13]

```
...  array([0.85984462, 0.14581783, 0.28219656, 0.79902342, 0.60618392])
```

```
np.random.rand(5,4)
```

[15]

```
...  array([[0.95404251, 0.62611004, 0.692988 , 0.53180799],
          [0.32360261, 0.58446612, 0.51718914, 0.32049731],
          [0.52852974, 0.7351838 , 0.63884682, 0.45631519],
          [0.04576068, 0.10476025, 0.04341695, 0.88704631],
          [0.77235022, 0.59780543, 0.80025621, 0.62047803]])
```

```
np.logspace(1,10, num=5,endpoint=True, base=10.0)
```

[19]

```
...  array([1.00000000e+01, 1.77827941e+03, 3.16227766e+05, 5.62341325e+07,
          1.00000000e+10])
```

```
a= np.array([[1,2,3],[4,5,6],[7,8,9]])  
a.shape  
[25]  
... (3, 3)
```

```
a[:,0]  
[45]  
... array([1, 4, 7])
```

```
a[0,:]  
[47]  
... array([1, 2, 3])
```

```
a[0,1]  
[41]  
... 2
```

```
a[1:3]  
[43]  
... array([[4, 5, 6],  
          [7, 8, 9]])
```

```
[51] a_sub=a[:2,:2]
      print(a_sub)
```

```
... [[1 2]
      [4 5]]
```

```
[57] a_sub[0,1]=10
      print(a_sub)
```

```
... [[ 1 10]
      [ 4  5]]
```

```
[93] s_col=np.append(s_row,[[17],[18],[19],[20]],axis=1)
      print(s_col)
```

```
... [[ 1 10  3 17]
      [ 4  5  6 18]
      [ 7  8  9 19]
      [10 11 14 20]]
```

```
[97] a_del=np.delete(s_col,1,axis=0)
      print(a_del)
```

```
... [[ 1 10  3 17]
      [ 7  8  9 19]
      [10 11 14 20]]
```

```
print(a)
```

```
[59]
```

```
... [[ 1 10  3]
      [ 4  5  6]
      [ 7  8  9]]
```

```
s_row=np.append(a,[[10,11,14]],axis=0)
print(s_row)
```

```
[81]
```

```
... [[ 1 10  3]
      [ 4  5  6]
      [ 7  8  9]
      [10 11 14]]
```

#Panda

import pandas as pd

```
[19] data1 = pd.read_csv('C:/Users/lab-2/mtcars.csv')
      print(data1)
```

```
...
      model  mpg  cyl  disp  hp  drat    wt    qsec  vs  am
0      Mazda RX4    21.0    6  160.0   110   3.90    2.620   16.46   0   1
1      Mazda RX4 Wag    21.0    6  160.0   110   3.90    2.875   17.02   0   1
2      Datsun 710    22.8    4  108.0    93   3.85    2.320   18.61   1   1
3      Hornet 4 Drive    21.4    6  258.0   110   3.08    3.215   19.44   1   0
4      Hornet Sportabout    18.7    8  360.0   175   3.15    3.440   17.02   0   0
5      Valiant    18.1    6  225.0   105   2.76    3.460   20.22   1   0
6      Duster 360    14.3    8  360.0   245   3.21    3.570   15.84   0   0
7      Merc 240D    24.4    4  146.7    62   3.69    3.190   20.00   1   0
8      Merc 230    22.8    4  140.8    95   3.92    3.150   22.90   1   0
9      Merc 280    19.2    6  167.6   123   3.92    3.440   18.30   1   0
10     Merc 280C    17.8    6  167.6   123   3.92    3.440   18.90   1   0
11     Merc 450SE    16.4    8  275.8   180   3.07    4.070   17.40   0   0
12     Merc 450SL    17.3    8  275.8   180   3.07    3.730   17.60   0   0
13     Merc 450SLC    15.2    8  275.8   180   3.07    3.780   18.00   0   0
14  Cadillac Fleetwood    10.4    8  472.0   205   2.93    5.250   17.98   0   0
15  Lincoln Continental    10.4    8  460.0   215   3.00    5.424   17.82   0   0
```

```
[21] data1.info()
```

```
...
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32 entries, 0 to 31
Data columns (total 12 columns):
 #   Column  Non-Null Count  Dtype
---  -
0   model   32 non-null     object
1   mpg     32 non-null     float64
2   cyl     32 non-null     int64
3   disp    32 non-null     float64
4   hp      32 non-null     int64
5   drat    32 non-null     float64
6   wt      32 non-null     float64
7   qsec    32 non-null     float64
8   vs      32 non-null     int64
9   am      32 non-null     int64
10  gear    32 non-null     int64
11  carb    32 non-null     int64
dtypes: float64(5), int64(6), object(1)
memory usage: 3.1+ KB
```

```
[27] data1.head()
```

```
...
```

	model	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2

```
data1.isnull().sum()
```

```
[33]
```

```
...    model      0
      mpg      0
      cyl      0
      disp      0
      hp       0
      drat      0
      wt       0
      qsec      0
      vs       0
      am       0
      gear      0
      carb      0
      dtype: int64
```

```
data1.isnull()
```

	model	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	False	False	False	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	False	False
5	False	False	False	False	False	False	False	False	False	False	False	False
6	False	False	False	False	False	False	False	False	False	False	False	False

```
data1.iat[8,4]
```

```
[47]
```

```
...    95
```

```
data1.shape
```

```
[37]
```

```
...    (32, 12)
```

```
data1.size
```

```
[39]
```

```
...    384
```



```
[41] data1.ndim
...    2

[45] data1.at[8,"model"]
...    'Merc 230'
```


```
[49] data1.loc[:, "model"]
...
0      Mazda RX4
1      Mazda RX4 Wag
2      Datsun 710
3      Hornet 4 Drive
4      Hornet Sportabout
5      Valiant
6      Duster 360
7      Merc 240D
8      Merc 230
9      Merc 280
10     Merc 280C
11     Merc 450SE
12     Merc 450SL
13     Merc 450SLC
14     Cadillac Fleetwood
15     Lincoln Continental
16     Chrysler Imperial
17     Fiat 128
18     Honda Civic
19     Toyota Corolla
20     Toyota Corona
21     Dodge Challenger
22     AMC Javelin
23     Camaro Z28
24     Pontiac Firebird
...
28     Ford Pantera L
29     Ferrari Dino
30     Maserati Bora
31     Volvo 142E
Name: model, dtype: object
```

Practical 5

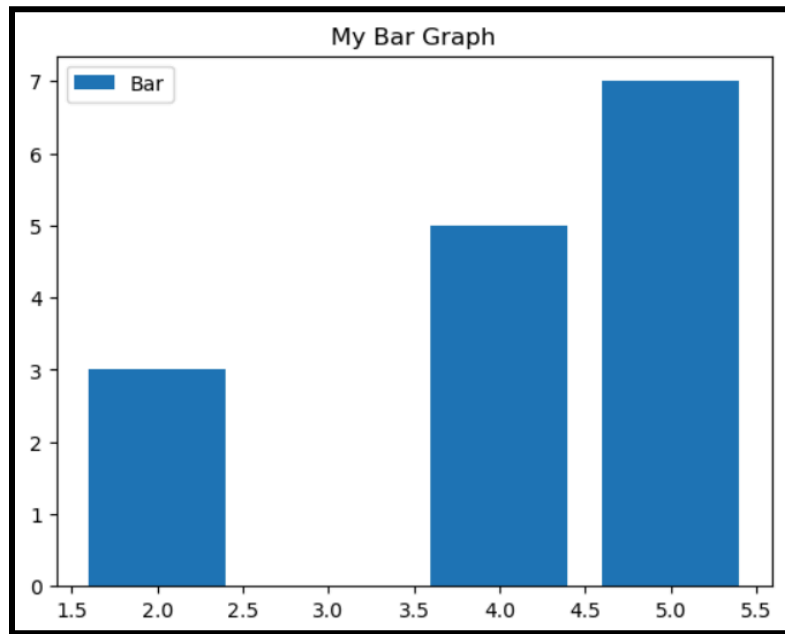
Aim: Intro to Python Libraries – Matplotlib, SciPy

```
import matplotlib.pyplot as plt
import pandas as pd


cars = pd.read_csv("C:/Users/USER/Downloads/mtcars.csv")
x = [2,4,5]
y = [3,5,7]
plt.plot(x,y)
plt.title("My Graph")
plt.legend(['Line'])
plt.show()
```



```
plt.bar(x,y)  
plt.title("My Bar Graph")  
plt.legend(["Bar"])  
plt.show()
```

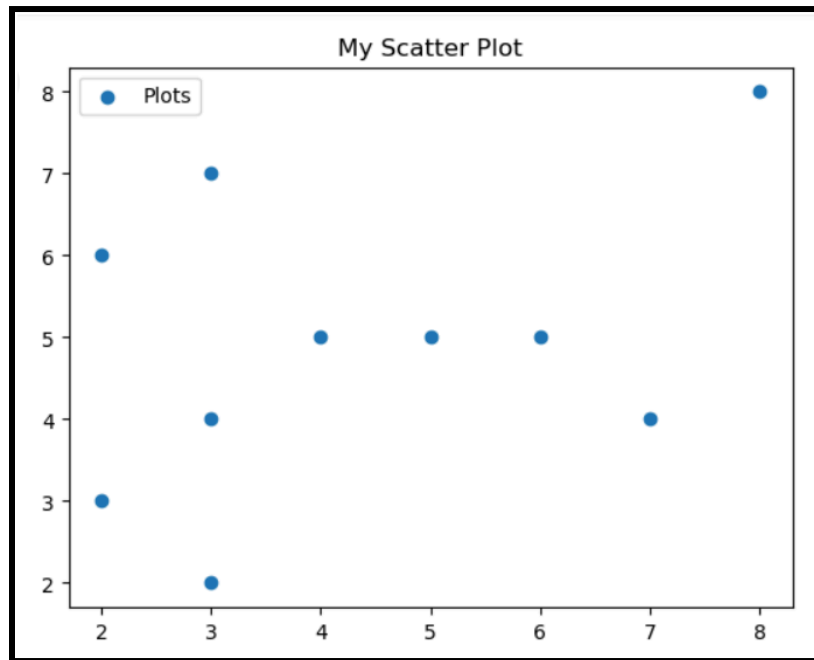


```
x = [2,3,4,5,2,3,6,7,8,4,5,6,0]  
plt.hist(x, bins = [1,2,3,4,5,6,7])  
plt.legend(['Bar'])  
plt.title("My Histogram")
```




```
x = [3,4,2,3,2,5,3,7,6,8]
y = [4,5,6,7,3,5,2,4,5,8]
```

```
plp.scatter(x,y)
plp.title("My Scatter Plot")
plp.legend(["Plots"])
plp.show()
```



```
labels = ['Mango', 'Banana', 'Watermelon', 'Strawberry']
sizes = [10,30,20,40]
plp.pie(cars['carb'],labels=cars['model'],autopct='%1.2f%%')
plp.show()
```

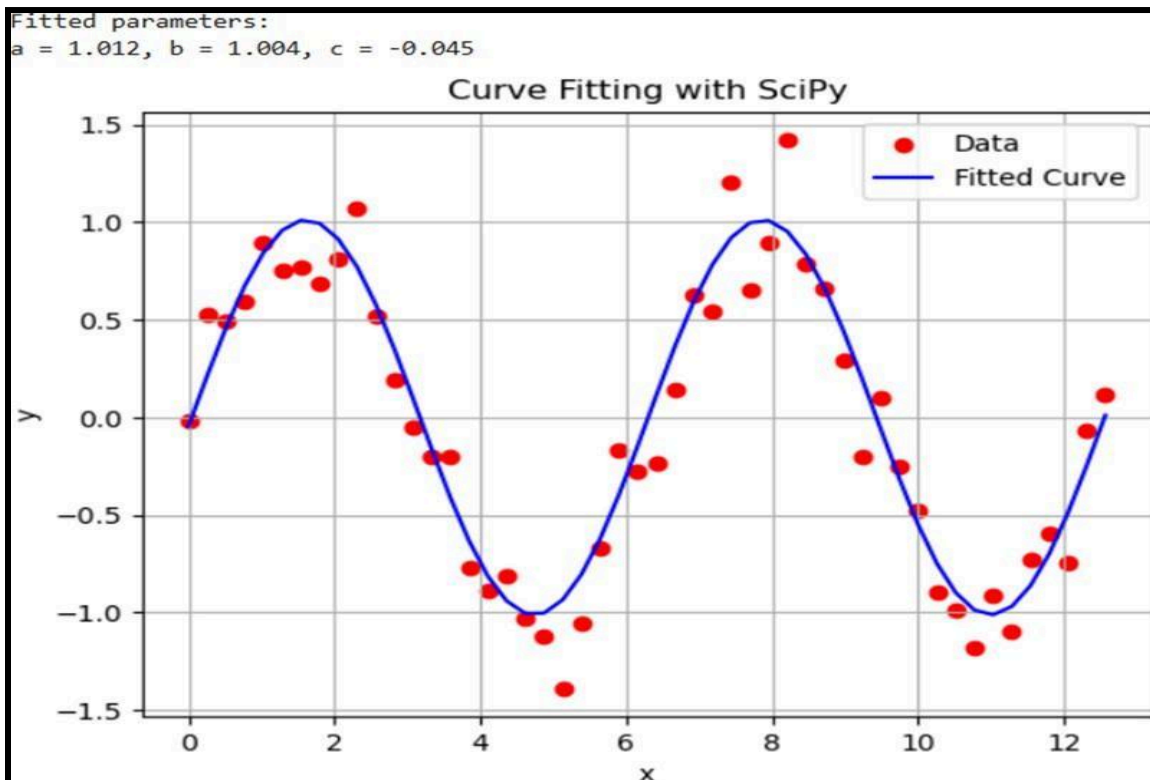


#scipy

```
import numpy as np
import matplotlib.pyplot as plt
from scipy.optimize import curve_fit
# Sample data (e.g., generated from a noisy sine wave) x_data = np.linspace(0, 4 * np.pi, 50)
y_data = np.sin(x_data) + 0.2 * np.random.normal(size=len(x_data)) # Define the model function to fit
def model_func(x, a, b, c): return a * np.sin(b * x + c)
# Fit the model to the data
params, params_covariance = curve_fit(model_func, x_data, y_data, p0=[1, 1, 0]) # Print the optimized parameters
print("Fitted parameters:")
print(f"a = {params[0]:.3f}, b = {params[1]:.3f}, c = {params[2]:.3f}")

# Plotting the data and the fitted curve plt.scatter(x_data, y_data, label='Data', color='red')
plt.plot(x_data, model_func(x_data, *params), label='Fitted Curve', color='blue') plt.legend()
plt.title('Curve Fitting with SciPy') plt.xlabel('x')
plt.ylabel('y') plt.grid(True) plt.show()
```

Output:



Practical 6

Aim: Intro to Python Libraries – Exploratory Data Analysis

Code:

```
import pandas as pd
from sklearn import metrics

df = pd.read_csv('C:/adeela/CreditRisk.csv')
print("DataFrame head:")
df.head()
```

Output

DataFrame head:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property_Area
0	LP001002	Male	No	0	Graduate	No	5849	0.0	0	360.0	1.0	Urban
1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128	360.0	1.0	Rural
2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66	360.0	1.0	Urban
3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	1.0	Urban
4	LP001008	Male	No	0	Graduate	No	6000	0.0	141	360.0	1.0	Urban

```
df.dtypes
num_cols = df.select_dtypes(include=np.number)
num_cols
```

Output:

```
Loan_ID      object
Gender       object
Married      object
Dependents   object
Education    object
Self_Employed object
ApplicantIncome  int64
CoapplicantIncome float64
LoanAmount     int64
Loan_Amount_Term float64
Credit_History float64
Property_Area  object
Loan_Status    int64
dtype: object
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status
0	5849	0.0	0	360.0	1.0	1
1	4583	1508.0	128	360.0	1.0	0
2	3000	0.0	66	360.0	1.0	1
3	2583	2358.0	120	360.0	1.0	1
4	6000	0.0	141	360.0	1.0	1
...
609	2900	0.0	71	360.0	1.0	1
610	4106	0.0	40	180.0	1.0	1
611	8072	240.0	253	360.0	1.0	1
612	7583	0.0	187	360.0	1.0	1
613	4583	0.0	133	360.0	0.0	0

614 rows × 6 columns

```
df['ApplicantIncome'].mean()
5403.459283387622
obj_cols=df.select_dtypes(exclude=['number']).columns
obj_cols

Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
      'Self_Employed', 'Property_Area'],
      dtype='object')
X = df.drop('Loan_Status', axis=1)
y =df['Loan_Status']
y.value_counts()
1      422
0      192
Name: Loan_Status, dtype: int64
pd.get_dummies(df,'Gender')
```

Output:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status	Gender_LP001002	Gender_LP001003	Gender_LP001005	Gender_LP00
0	5549	0.0	0	360.0	1.0	1	1	0	0	
1	4583	1508.0	128	360.0	1.0	0	0	1	0	
2	3000	0.0	66	360.0	1.0	1	0	0	1	
3	2583	2358.0	120	360.0	1.0	1	0	0	0	
4	6000	0.0	141	360.0	1.0	1	0	0	0	
--	--	--	--	--	--	--	--	--	--	--
609	2900	0.0	71	360.0	1.0	1	0	0	0	
610	4106	0.0	40	180.0	1.0	1	0	0	0	
611	8072	240.0	253	360.0	1.0	1	0	0	0	
612	7583	0.0	187	360.0	1.0	1	0	0	0	
613	4583	0.0	133	360.0	0.0	0	0	0	0	

614 rows x 11 columns

```
correlation_matrix =num_cols.corr()

correlation_matrix
```


Output:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status
ApplicantIncome	1.000000	-0.116605	0.538290	-0.045306	-0.014715	-0.004710
CoapplicantIncome	-0.116605	1.000000	0.190377	-0.059878	-0.002056	-0.059187
LoanAmount	0.538290	0.190377	1.000000	0.040539	-0.002197	-0.010631
Loan_Amount_Term	-0.045306	-0.059878	0.040539	1.000000	0.001470	-0.021268
Credit_History	-0.014715	-0.002056	-0.002197	0.001470	1.000000	0.561678
Loan_Status	-0.004710	-0.059187	-0.010631	-0.021268	0.561678	1.000000

```
#Fill Null Values
100 * credit_df.isnull().sum() / credit_df.shape[0]
```

Output:

```
Loan_ID      0.000000
Gender       2.117264
Married      0.488599
Dependents   2.442997
Education    0.000000
Self_Employed 5.211726
ApplicantIncome 0.000000
CoapplicantIncome 0.000000
LoanAmount   0.000000
Loan_Amount_Term 2.280130
Credit_History 8.143322
Property_Area 0.000000
Loan_Status  0.000000
dtype: float64
```

```
DF=credit_df.drop(credit_df.columns[0],axis=1)
DF.head()
```

Output:

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property_Area	Loan_Status
0	Male	No	0	Graduate	No	5849	0.0	0	360.0	1.0	Urban	
1	Male	Yes	1	Graduate	No	4583	1508.0	128	360.0	1.0	Rural	
2	Male	Yes	0	Graduate	Yes	3000	0.0	66	360.0	1.0	Urban	
3	Male	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	1.0	Urban	
4	Male	No	0	Graduate	No	6000	0.0	141	360.0	1.0	Urban	

```
object_columns = DF.select_dtypes(include=['object']).columns
numeric_columns = DF.select_dtypes(exclude=['object']).columns
#credit_df.columns[credit_df.dtypes == object]
#credit_df.columns[credit_df.dtypes == object]
for column in object_columns:
    majority = DF[column].value_counts().iloc[0]
    DF[column].fillna(majority, inplace=True)
for column in numeric_columns:
    mean = DF[column].mean()
    DF[column].fillna(mean, inplace=True)
# Impute
DF.head()
```

Output:

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property_Area	Loan_Status
0	Male	No	0	Graduate	No	5849	0.0	0	360.0	1.0	Urban	
1	Male	Yes	1	Graduate	No	4583	1508.0	128	360.0	1.0	Rural	
2	Male	Yes	0	Graduate	Yes	3000	0.0	66	360.0	1.0	Urban	
3	Male	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	1.0	Urban	
4	Male	No	0	Graduate	No	6000	0.0	141	360.0	1.0	Urban	

Practical 7

Aim: Implementation of Perceptron Algorithm

Code: perceptron.ipynb

```
import numpy as np

def perceptron_or(x1, x2): w1 = 1
w2 = 1
b = -0.5
result = w1 * x1 + w2 * x2 + b

if result >= 0: return 1
else:
return 0

print("Output: ")
print(perceptron_or(0, 0))
print(perceptron_or(0, 1))
print(perceptron_or(1, 0))
print(perceptron_or(1, 1))
```

Output:

0

1

1

1

Practical 8

Aim: Implementation of Adaline for AND Operations

Code: adaline_and_operation.ipynb

```
import numpy as np

class Adaline:
    def __init__(self, learning_rate=0.01, n_iter=100):
        self.learning_rate = learning_rate
        self.n_iter = n_iter
        self.weights = None
        self.bias = None

    def fit(self, X, y):
        # Initialize weights and bias
        self.weights = np.zeros(X.shape[1])
        self.bias = 0

        # Perform gradient descent
        for _ in range(self.n_iter):
            # Calculate net input (weighted sum of inputs)
            net_input = np.dot(X, self.weights) + self.bias

            # Calculate error (difference between prediction and actual)
            error = y - net_input

            # Update weights and bias using gradient descent
            self.weights += self.learning_rate * np.dot(X.T, error)
            self.bias += self.learning_rate * np.sum(error)

    def predict(self, X):
        # Calculate net input
        net_input = np.dot(X, self.weights) + self.bias

        # Apply a threshold (0.0 for linear activation)
        return np.where(net_input >= 0.0, 1, 0)

# Example Usage (AND operation)
X = np.array([[0, 0], [0, 1], [1, 0], [1, 1]])
y = np.array([0, 0, 0, 1])

# Create and train Adaline model
ada = Adaline(learning_rate=0.1, n_iter=1000)
ada.fit(X, y)

# Make predictions
predictions = ada.predict(X)
print("Output: ")
print("Predictions:", predictions)
```

Output:

Predictions: [0 1 1 1]

Practical 9

Aim: Implementation of Gradient Descent Algorithm

Code: gradient_descent.ipynb

```
def predict(row, weights):  
    activation = weights[0]  
    for i in range(len(row)-1):  
        activation += weights[i + 1] * row[i] return 1.0 if activation >= 0.0 else 0.0  
  
# test predictions  
dataset = [  
    [2.7810836, 2.550537003, 0],  
    [1.465489372, 2.362125076, 0],  
    [3.396561688, 4.400293529, 0],  
    [1.38807019, 1.850220317, 0],  
    [3.06407232, 3.005305973, 0],  
    [7.627531214, 2.759262235, 1],  
    [5.332441248, 2.088626775, 1],  
    [6.922596716, 1.77106367, 1],  
    [8.675418651, -0.242068655, 1],  
    [7.673756466, 3.508563011, 1]  
]  
  
weights = [-0.1, 0.20653640140000007, -0.23418117710000003]  
  
for row in dataset:  
    prediction = predict(row, weights)  
    print("Expected=%d, Predicted=%d" % (row[-1], prediction))
```

Output:

```
Expected=0, Predicted=0
Expected=0, Predicted=0
Expected=0, Predicted=0
Expected=0, Predicted=0
Expected=0, Predicted=0
Expected=1, Predicted=1
Expected=1, Predicted=1
Expected=1, Predicted=1
Expected=1, Predicted=1
Expected=1, Predicted=1
```

```
# Function to make predictions using weights def
predict(row, weights):
activation = weights[0]
for i in range(len(row) - 1):
    activation += weights[i + 1] * row[i] return 1.0 if
activation >= 0.0 else 0.0

# Function to train perceptron weights using stochastic gradient descent def
train_weights(train, l_rate, n_epoch):
    weights = [0.0 for i in range(len(train[0]))] # Initialize weights to 0.0 for each feature
    for epoch in range(n_epoch): sum_error
        = 0.0

    for row in train:
        prediction = predict(row, weights)    # Make prediction using current
        weights
        error = row[-1] - prediction    # Calculate error as actual -
        predicted
        sum_error += error ** 2    # Accumulate squared error
        weights[0] = weights[0] + l_rate * error    # Update bias
        (weights[0])
    for i in range(len(row) - 1):
        weights[i + 1] = weights[i + 1] + l_rate * error * row[i] #
    Update weights for features
    print('>epoch=%d, lrate=%.3f, error=%.3f' % (epoch, l_rate, sum_error))

# Print epoch details
    return weights # Return trained weights
# Dataset for training
dataset = [
    [2.7810836, 2.550537003, 0],
    [1.465489372, 2.362125076, 0],
    [3.396561688, 4.400293529, 0],
    [1.38807019, 1.850220317, 0],
    [3.06407232, 3.005305973, 0],
    [7.627531214, 2.759262235, 1],
```

```
[5.332441248, 2.088626775, 1],  
[6.922596716, 1.77106367, 1],  
[8.675418651, -0.242068655, 1],  
[7.673756466, 3.508563011, 1]  
]
```

```
l_rate = 0.1 # Learning rate  
n_epoch = 5 # Number of epochs for training  
# Train weights using the dataset  
weights = train_weights(dataset, l_rate, n_epoch)  
print(weights)  
# Print the trained weights
```

Output:

```
>epoch=0, lrate=0.100, error=2.000  
>epoch=1, lrate=0.100, error=1.000  
>epoch=2, lrate=0.100, error=0.000  
>epoch=3, lrate=0.100, error=0.000  
>epoch=4, lrate=0.100, error=0.000  
[-0.1, 0.20653640140000007, -0.23418117710000003]
```

Practical 10

Aim: Implementation of Principal Component Analysis

Code:

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
```

#import the Datasets

```
data = pd.read_csv("driver-data.csv")
print(data.head())
print(data.info())
```

```
      id  mean_dist_day  mean_over_speed_perc
0  3423311935         71.24                 28
1  3423313212         52.53                 25
2  3423313724         64.54                 27
3  3423311373         55.69                 22
4  3423310999         54.58                 25
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4000 entries, 0 to 3999
Data columns (total 3 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    4000 non-null   int64
1   mean_dist_day         4000 non-null   float64
2   mean_over_speed_perc  4000 non-null   int64
dtypes: float64(1), int64(2)
memory usage: 93.9 KB
None
```

#Feature Scaling

```
x = data.values
print(x)
x = StandardScaler().fit_transform(x)
print(x)
```

```
[[ 3.42331194e+09  7.12400000e+01  2.80000000e+01]
 [ 3.42331321e+09  5.25300000e+01  2.50000000e+01]
 [ 3.42331372e+09  6.45400000e+01  2.70000000e+01]
 ...
 [ 3.42331292e+09  1.70910000e+02  1.20000000e+01]
 [ 3.42331363e+09  1.76140000e+02  5.00000000e+00]
 [ 3.42331153e+09  1.68030000e+02  9.00000000e+00]]
[[ -0.44383803 -0.0898104   1.26061251]
 [  0.66207644 -0.43977285  1.04174351]
 [  1.10548146 -0.215131   1.18765617]
```


#Creating an instance of the PCA class with component=2

```
pca = PCA(n_components=2)
principle_component = pca.fit_transform(x)
```


#Creating a new dataset for PCA

```
principle_data = pd.DataFrame(data=principle_component, columns=['(Principle Component 1)', '(Principle Component 2)'])
print(principle_data)
```

	(Principle Component 1)	(Principle Component 2)
0	0.834991	-0.317476
1	0.412104	0.784128
2	0.665867	1.223204
3	0.329080	-0.822333
4	0.474999	-1.122316
...
3995	1.103304	-1.626122
3996	1.028275	-0.032743
3997	1.314121	0.297566
3998	1.011236	0.853422
3999	1.143851	-0.913348

[4000 rows x 2 columns]

```
plt.figure(figsize=(8,9))
plt.scatter(principle_data['(Principle Component 1)'], principle_data['(Principle Component 2)'])
plt.xlabel('Principle Component 1')
plt.ylabel('Principle Component 2')
plt.title('PCA Visualization')
plt.show()
```



```
# Explained variance ratio
explained_variance_ratio = pca.explained_variance_ratio_
print("Explained Variance Ratio:", explained_variance_ratio)
print("Total Explained Variance:", np.sum(explained_variance_ratio))
```

```
Explained Variance Ratio: [0.4223297  0.33447006]
Total Explained Variance: 0.7567997587537549
```

Practical 11

Aim: Implementation of Normalization and Transformation

```
from sklearn import preprocessing import
numpy as np
x_array = np.array([2,3,5,6,7,4,8,7,6])
normalized_arr = preprocessing.normalize([x_array]) print(normalized_arr)
```

```
[[ 0.11785113  0.1767767  0.29462783  0.35355339  0.41247896  0.23570226
  0.47140452  0.41247896  0.35355339]]
```

```
import pandas as pd import
numpy as np

# Sample data
data = {'Name': ['Alice', 'Bob', 'Charlie', 'David'], 'Age': [25, 30, 22,
28],
        'City': ['New York', 'London', 'Paris', 'Tokyo']} df =
pd.DataFrame(data)

print("Original DataFrame:") print(df)
```

Transformation

```
Original DataFrame:
   Name  Age  City
0  Alice   25 New York
1   Bob   30  London
2 Charlie   22   Paris
3  David   28   Tokyo
```

```
#Adding a new column
df['Age_Group'] = pd.cut(df['Age'], bins=[18, 25, 30, 100], labels=['Young', 'Adult', 'Senior'])
print("\nDataFrame with Age Group column:")
print(df)
```

```
DataFrame with Age Group column:
   Name  Age  City Age_Group
0  Alice   25 New York   Young
1   Bob   30  London   Adult
2 Charlie   22   Paris   Young
3  David   28   Tokyo   Adult
```

```
# 2. Creating dummy variables for categorical features df =
pd.get_dummies(df, columns=['City'])
print("\nDataFrame after creating dummy variables for City:") print(df)
```

```
DataFrame after creating dummy variables for City:
   Name  Age Age_Group  City_London  City_New York  City_Paris  City_Tokyo
0  Alice   25   Young         False             True         False         False
1   Bob   30   Adult          True             False         False         False
2 Charlie   22   Young         False             False         True         False
3  David   28   Adult         False             False         False         True
```

```
import pandas as pd
from sklearn import metrics

df = pd.read_csv('CreditRisk.csv')
print("DataFrame head:")
df.head()
```

```
DataFrame head:
   Loan_ID  Gender  Married  Dependents  Education  Self_Employed  ApplicantIncome  CoapplicantIncome  LoanAmount  Loan_Amount_Term  Credit_History  Property_Area  Loan_Status
0  LP001002   Male     No         0      Graduate         No           5849              0.0           0           360.0           1.0      Urban         1
1  LP001003   Male     Yes        1      Graduate         No           4583             1508.0         128           360.0           1.0      Rural         0
2  LP001005   Male     Yes        0      Graduate         Yes           3000              0.0           66           360.0           1.0      Urban         1
3  LP001006   Male     Yes        0  Not Graduate         No           2583             2358.0         120           360.0           1.0      Urban         1
4  LP001008   Male     No         0      Graduate         No           6000              0.0          141           360.0           1.0      Urban         1
```

```
df.dtypes
```

Output:

```
Loan_ID          object
Gender           object
Married          object
Dependents       object
Education        object
Self_Employed    object
ApplicantIncome  int64
CoapplicantIncome float64
LoanAmount       int64
Loan_Amount_Term float64
Credit_History   float64
Property_Area    object
Loan_Status      int64
dtype: object
```

```
#Example feature extraction:
# 1. Calculate the mean of a numerical column
#num_cols=df.select_dtypes(include=['number']).columns
num_cols = df.select_dtypes(include=np.number)
num_cols
```

Output:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status
0	5849	0.0	0	360.0	1.0	1
1	4583	1508.0	128	360.0	1.0	0
2	3000	0.0	66	360.0	1.0	1
3	2583	2358.0	120	360.0	1.0	1
4	6000	0.0	141	360.0	1.0	1
...
609	2900	0.0	71	360.0	1.0	1
610	4106	0.0	40	180.0	1.0	1
611	8072	240.0	253	360.0	1.0	1
612	7583	0.0	187	360.0	1.0	1
613	4583	0.0	133	360.0	0.0	0

614 rows × 6 columns

```
df['ApplicantIncome'].mean()
```

Output:

```
np.float64(5403.459283387622)
```

```
obj_cols=df.select_dtypes(exclude=['number']).columns
obj_cols
```

Output:

```
Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
       'Self_Employed', 'Property_Area'],
      dtype='object')
```

```
X = df.drop('Loan_Status', axis=1)
y =df['Loan_Status']
y.value_counts()
```

Output:

```
Loan_Status
1    422
0    192
Name: count, dtype: int64
```

```
pd.get_dummies(df, 'Gender')
```

Output:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status	Gender_LP001002	Gender_LP001003	Gender_LP001005	Gender_LP001006
0	5849	0.0	0	360.0	1.0	1	True	False	False	False
1	4583	1508.0	128	360.0	1.0	0	False	True	False	False
2	3000	0.0	66	360.0	1.0	1	False	False	True	False
3	2583	2358.0	120	360.0	1.0	1	False	False	False	True
4	6000	0.0	141	360.0	1.0	1	False	False	False	False
...
609	2900	0.0	71	360.0	1.0	1	False	False	False	False
610	4106	0.0	40	180.0	1.0	1	False	False	False	False
611	8072	240.0	253	360.0	1.0	1	False	False	False	False
612	7583	0.0	187	360.0	1.0	1	False	False	False	False
613	4583	0.0	133	360.0	0.0	0	False	False	False	False

614 rows × 11 columns

```
correlation_matrix =num_cols.corr()
correlation_matrix
```

Output:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status
ApplicantIncome	1.000000	-0.116605	0.538290	-0.045306	-0.014715	-0.004710
CoapplicantIncome	-0.116605	1.000000	0.190377	-0.059878	-0.002056	-0.059187
LoanAmount	0.538290	0.190377	1.000000	0.040539	-0.002197	-0.010631
Loan_Amount_Term	-0.045306	-0.059878	0.040539	1.000000	0.001470	-0.021268
Credit_History	-0.014715	-0.002056	-0.002197	0.001470	1.000000	0.561678
Loan_Status	-0.004710	-0.059187	-0.010631	-0.021268	0.561678	1.000000

Practical 12

Aim: Implementation of Logistic Regression

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sb

data = pd.read_csv("/content/CreditRisk.csv")
print(data)
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	
..	
609	LP002978	Female	No	0	Graduate	No	
610	LP002979	Male	Yes	3+	Graduate	No	
611	LP002983	Male	Yes	1	Graduate	No	
612	LP002984	Male	Yes	2	Graduate	No	
613	LP002990	Female	No	0	Graduate	Yes	

```
data.shape
```

```
(614, 13)
```


```
data.Loan_Status.value_counts()
```

count	
Loan_Status	
1	422
0	192

```
data.groupby(['Education','Loan_Status']).Education.count()
```

Education		
Education	Loan_Status	
Graduate	0	140
	1	340
Not Graduate	0	52
	1	82

```
sb.barplot(y = "Credit_History",x = "Loan_Status",hue = "Education",data = data)
```



#Finding the Null Values

```
data.isnull().sum()
```

	0
Loan_ID	0
Gender	13
Married	3
Dependents	15
Education	0
Self_Employed	32

ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	0
Loan_Amount_Term	14
Credit_History	50
Property_Area	0
Loan_Status	0
dtype: int64	

```
data = data.drop(data.columns[0],axis=1)
```

```
data.head()
```

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property_Area	Loan_Status
0	Male	No	0	Graduate	No	5849	0.0	0	360.0	1.0	Urban	1
1	Male	Yes	1	Graduate	No	4583	1508.0	128	360.0	1.0	Rural	0
2	Male	Yes	0	Graduate	Yes	3000	0.0	66	360.0	1.0	Urban	1
3	Male	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	1.0	Urban	1
4	Male	No	0	Graduate	No	6000	0.0	141	360.0	1.0	Urban	1

#Segregetting the columns into the different category

```
data_object = data.select_dtypes(include=['object']).columns
data_numeric = data.select_dtypes(exclude=['object']).columns
print(data_object)
print(data_numeric)
```



```
Index(['Gender', 'Married', 'Dependents', 'Education', 'Self_Employed',
      'Property_Area'],
      dtype='object')
Index(['ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
      'Loan_Amount_Term', 'Credit_History', 'Loan_Status'],
      dtype='object')
```

#Impute

for columns in data_object:

```
    majority = data[columns].value_counts().iloc[0]
    data.fillna(majority, inplace = True)
```

for columns in data_numeric:

```
    majority = data[columns].value_counts().iloc[0]
    data.fillna(majority, inplace = True)
```

data.isnull().sum()

	0
Gender	0
Married	0
Dependents	0
Education	0
Self_Employed	0
ApplicantIncome	0

CoapplicantIncome	0
LoanAmount	0
Loan_Amount_Term	0
Credit_History	0
Property_Area	0
Loan_Status	0
dtype: int64	

#Changing the String Values to Dummy Integer Values

```
df_dummy = pd.get_dummies(data, columns = data_object)
```

```
df_dummy.head()
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status	Gender_489	Gender_Female	Gender_Male	Married_489	...	Dependents_2	Dependents_3+
0	5849	0.0	0	360.0	1.0	1	False	False	True	False	...	False	False
1	4583	1508.0	128	360.0	1.0	0	False	False	True	False	...	False	False
2	3000	0.0	66	360.0	1.0	1	False	False	True	False	...	False	False
3	2583	2358.0	120	360.0	1.0	1	False	False	True	False	...	False	False
4	6000	0.0	141	360.0	1.0	1	False	False	True	False	...	False	False

5 rows × 25 columns


#Model Training

```
from sklearn.model_selection import train_test_split # For splitting the dataset into training and testing sets
from sklearn.linear_model import LogisticRegression # For logistic regression model
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report # For evaluating performance
```

```
X = df_dummy.drop('Loan_Status', axis = 1)
y = df_dummy.Loan_Status
```

```
train_x, test_x, train_y, test_y = train_test_split(X, y, test_size = 0.3, random_state=42)
```

```
train_x.shape, test_x.shape
model = LogisticRegression(max_iter=1000)
model.fit(train_x, train_y)
```



The screenshot shows a Jupyter Notebook cell with a blue header bar containing the text 'LogisticRegression' and two circular icons (one with an 'i' and one with a '?'). Below the header bar, the text 'LogisticRegression(max_iter=1000)' is displayed in a monospaced font.

```
y_pred = model.predict(test_x)
```

```
# Calculate accuracy
```


```
accuracy = accuracy_score(test_y, y_pred)
print(f"Accuracy: {accuracy}")
```

```
# Generate the confusion matrix
```

```
conf_matrix = confusion_matrix(test_y, y_pred)
```

```
print("Confusion Matrix:")
```

```
print(conf_matrix)
```



The screenshot shows the output of the Python code in a monospaced font. It displays the accuracy value '0.6216216216216216' and the confusion matrix as a 2x2 array: '[[3 62]' and '[8 112]]'.

```
# Generate the classification report
```

```
class_report = classification_report(test_y, y_pred)
```

```
print("Classification Report:")
```

```
print(class_report)
```

Classification Report:

	precision	recall	f1-score	support
0	0.27	0.05	0.08	65
1	0.64	0.93	0.76	120
accuracy			0.62	185
macro avg	0.46	0.49	0.42	185
weighted avg	0.51	0.62	0.52	185

Practical 13

Aim: Implementation of Support Vector Machine – RBF Kernel

Code: (Here we are considering same CreditRisk.csv)


#Model Training and Data Splitting

```
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

X = df_dummy.drop('Loan_Status', axis = 1)
y = df_dummy.Loan_Status

train_x, test_x, train_y, test_y = train_test_split(X, y, test_size = 0.3, random_state=42)

train_x.shape, test_x.shape
svm_model = SVC(kernel='rbf', gamma=0.00001, C=1000)
svm_model.fit(train_x, train_y)
```



#Prediction

```
train_y_hat = svm_model.predict(train_x)
test_y_hat = svm_model.predict(test_x)

print('-'*20, 'Train', '-'*20)

print(classification_report(train_y, train_y_hat))
print('-'*20, 'Test', '-'*20)
print(classification_report(test_y, test_y_hat))
print(accuracy_score(test_y, test_y_hat))

confusion_matrix(test_y, test_y_hat)
```

```

----- Train -----
      precision    recall  f1-score   support

     0       0.97       0.95       0.96         127
     1       0.98       0.99       0.98         302

 accuracy          0.98         429
 macro avg       0.97       0.97       0.97         429
weighted avg       0.98       0.98       0.98         429

----- Test -----
      precision    recall  f1-score   support

     0       0.33       0.22       0.26          65
     1       0.64       0.77       0.70         120

 accuracy          0.57         185
 macro avg       0.49       0.49       0.48         185
weighted avg       0.53       0.57       0.55         185

0.572972972972973

```

```

array([[14, 51],
       [28, 92]])

```

Practical 14

Aim: Implementing Elbow Method for Choosing Number of Clusters

Code:

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler

df = pd.read_csv('driver-data.csv')
print(df.info())
print(df.shape)

data = df[['mean_dist_day', 'mean_over_speed_perc']]
print(data)

data.fillna(data.mean(), inplace=True)
print(data)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4000 entries, 0 to 3999
Data columns (total 3 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   id                    4000 non-null   int64   
 1   mean_dist_day         4000 non-null   float64  
 2   mean_over_speed_perc  4000 non-null   int64   
dtypes: float64(1), int64(2)
memory usage: 93.9 KB
None
(4000, 3)
   mean_dist_day  mean_over_speed_perc
0          71.24                28
1          52.53                25
2          64.54                27
3          55.69                22
4          54.58                25
...          ...                  ...
3995         160.04                10
3996         176.17                 5
3997         170.91                12
3998         176.14                 5
3999         168.03                 9
```

```
[4000 rows x 2 columns]
   mean_dist_day  mean_over_speed_perc
0          71.24                28
1          52.53                25
2          64.54                27
3          55.69                22
4          54.58                25
...          ...                  ...
3995         160.04                10
3996         176.17                 5
3997         170.91                12
3998         176.14                 5
3999         168.03                 9

[4000 rows x 2 columns]
```


```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(data)

wcss = []
for i in range(1,16):
    kmeans = KMeans(n_clusters=i, init='k-means++', random_state=42, n_init=10)
```

```
kmeans.fit(X_scaled)  
wcss.append(kmeans.inertia_)  
print(wcss)
```

```
[8000.0, 3911.9263904284157, 1756.5445821314272, 739.153450864558,  
619.4037594867996, 502.03685490351063, 437.8780702842313, 374.8024199852508,  
337.1456744336048, 306.97337203496494, 276.4393166945794, 248.09311615011552,  
229.74448771226085, 210.66993580924535, 198.0806841488063]
```

```
# Plot the Elbow method graph  
plt.plot(range(1, 16), wcss)  
plt.title('Elbow Method')  
plt.xlabel('Number of Clusters')  
plt.ylabel('WCSS')  
plt.show()
```



Hence 3 Number of Clusters would be great.

Practical 15

Aim: Ensemble Techniques – Bagging, Boosting, Stacking, Voting

Code:

```
from sklearn.ensemble import BaggingClassifier, AdaBoostClassifier, StackingClassifier,
VotingClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.datasets import load_iris
from sklearn.metrics import accuracy_score
```

```
iris = load_iris()
X, y = iris.data, iris.target

X.shape
```

```
(150, 4)
```

```
y.shape
```

```
(150,)
```

```
# Splitting Data into training and testing
x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# 1. Bagging (using Decision Tree as base estimator)
bagging_model = BaggingClassifier(estimator=DecisionTreeClassifier(), n_estimators=10,
random_state=42)
bagging_model.fit(x_train, y_train)
bagging_predictions = bagging_model.predict(x_test)
bagging_accuracy = accuracy_score(y_test, bagging_predictions)
print("Bagging Accuracy:", bagging_accuracy)
```

```
Baggin Accuracy: 1.0
```



```
# 2. Boosting (using AdaBoost with Decision Tree)
boosting_model = AdaBoostClassifier(estimator=DecisionTreeClassifier(), n_estimators=50,
random_state=42)
boosting_model.fit(x_train, y_train)
boosting_predictions = boosting_model.predict(x_test)
boosting_accuracy = accuracy_score(y_test, boosting_predictions)
print("Boosting Accuracy:", boosting_accuracy)
```

Boosting Accuracy: 1.0

```
estimators = [
    ('dt', DecisionTreeClassifier()),
    ('lr', LogisticRegression()),
    ('knn', KNeighborsClassifier())
]
```

```
# 3. Stacking (using Decision Tree, logistic Regression, and KNN as base estimators)
stacking_model = StackingClassifier(estimators=estimators,
final_estimator=LogisticRegression())
stacking_model.fit(x_train, y_train)
stacking_predictions = stacking_model.predict(x_test)
stacking_accuracy = accuracy_score(y_test, stacking_predictions)
print("Stacking Accuracy:", stacking_accuracy)
```

Stacking Accuracy: 1.0

```
# Voting (using Decision Tree, Logistic Regression and KNN)
voting_model = VotingClassifier(estimators=estimators, voting='hard')
voting_model.fit(x_train, y_train)
voting_predictions = voting_model.predict(x_test)
voting_accuracy = accuracy_score(y_test, voting_predictions)
print("Voting Accuracy:", voting_accuracy)
```

Voting Accuracy: 1.0

Practical 16

Aim: Implementing Bagging and Voting Algorithm Using Random Forest as Base Estimator

Code:

```
# Import necessary libraries
import numpy as np
import pandas as pd
from sklearn.metrics import accuracy_score
from sklearn.ensemble import RandomForestClassifier, BaggingClassifier, VotingClassifier
from sklearn.model_selection import train_test_split
```

```
# Load the Titanic dataset
train = pd.read_csv("titanic.csv") print(train.shape) # Output:
(891, 12)
```

```
(891, 12)
```

```
# Checking for missing data
NAs = pd.concat([train.isnull().sum()], axis=1, keys=["Train"])
NAs[NAs.sum(axis=1) > 0] # Display columns with missing values
```

	Train
Age	177
Cabin	687
Embarked	2

```
# Removing unnecessary columns
train.pop("Cabin")
train.pop("Name")
train.pop("Ticket")
```

```
0          A/5 21171
1          PC 17599
2      STON/O2. 3101282
3          113803
4          373450
...
886          211536
887          112053
888      W./C. 6607
889          111369
890          370376
Name: Ticket, Length: 891, dtype: object
```

```
# Filling missing Age values with mean
train['Age'] = train['Age'].fillna(train['Age'].mean())

# Filling missing Embarked values with most common value
train['Embarked'] = train['Embarked'].fillna(train['Embarked'].mode()[0])

# Converting Pclass to string
train['Pclass'] = train['Pclass'].apply(str)

# Getting Dummies for categorical variables
for col in train.dtypes[train.dtypes == "object"].index:
    for_dumy = train.pop(col)
    train = pd.concat([train, pd.get_dummies(for_dumy, prefix=col)], axis=1)

train.head()
```

PassengerId	Survived	Age	SibSp	Parch	Fare	Pclass_1	Pclass_2	Pclass_3	Sex_female	Sex_male	Embarked_C	Embarked_Q	Embarked_S
0	1	0 22.0	1	0	7.2500	False	False	True	False	True	False	False	True
1	2	1 38.0	1	0	71.2833	True	False	False	True	False	True	False	False
2	3	1 26.0	0	0	7.9250	False	False	True	True	False	False	False	True
3	4	1 35.0	1	0	53.1000	True	False	False	True	False	False	False	True
4	5	0 35.0	0	0	8.0500	False	False	True	False	True	False	False	True

```
# Separating labels
labels = train.pop("Survived")

# Splitting the dataset into training and testing sets
x_train, x_test, y_train, y_test = train_test_split(train, labels, test_size=0.25)

# Importing and training RandomForestClassifier
rf = RandomForestClassifier(n_estimators=10)
rf.fit(x_train, y_train)
```

▼ RandomForestClassifier ⓘ ?

```
RandomForestClassifier(n_estimators=10)
```

```
# Making predictions
y_pred = rf.predict(x_test)

# Accuracy
accuracy = accuracy_score(y_test, y_pred)
accuracy
```

0.8071748878923767

```
# 1. Bagging (using RandomForestClassifier as base estimator)
bagging_model = BaggingClassifier(estimator=RandomForestClassifier(), n_estimators=10)
bagging_model.fit(x_train, y_train)
bagging_prediction = bagging_model.predict(x_test)
bagging_accuracy = accuracy_score(y_test, bagging_prediction)
print("Bagging Accuracy", bagging_accuracy)
```

Bagging Accuracy 0.8026905829596412

```
# 4. Voting (using RandomForestClassifier)
voting_model = VotingClassifier(estimators=[
    ('rf', RandomForestClassifier())
], voting='hard') # 'hard' for majority vote, 'soft' for weighted average probabilities
voting_model.fit(x_train, y_train)
voting_predictions = voting_model.predict(x_test)
voting_accuracy = accuracy_score(y_test, voting_predictions)
print("Voting Accuracy:", voting_accuracy)
```

Voting Accuracy: 0.7982062780269058

Practical 17

Aim: Implementing AdaBoost Algorithm

Code:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report
```

```
train = pd.read_csv('titanic.csv') print(train)
```

	PassengerId	Survived	Pclass	...	Fare	Cabin	Embarked
0	1	0	3	...	7.2500	NaN	S
1	2	1	1	...	71.2833	C85	C
2	3	1	3	...	7.9250	NaN	S
3	4	1	1	...	53.1000	C123	S
4	5	0	3	...	8.0500	NaN	S
..
886	887	0	2	...	13.0000	NaN	S
887	888	1	1	...	30.0000	B42	S
888	889	0	3	...	23.4500	NaN	S
889	890	1	1	...	30.0000	C148	C
890	891	0	3	...	7.7500	NaN	Q

[891 rows x 12 columns]

```
# Checking for missing data
NAs = pd.concat([train.isnull().sum()], axis=1, keys=["Train"])
NAs[NAs.sum(axis=1) > 0]# Display columns with missing values
```

Train	
Age	177
Cabin	687
Embarked	2

```
# Removing unnecessary columns
train.pop("Cabin")
train.pop("Name")
train.pop("Ticket")
```

```
0      A/5  21171
1      PC  17599
2  STON/O2. 3101282
3      113803
4      373450
...
886      211536
887      112053
888  W./C.  6607
889      111369
890      370376
Name: Ticket, Length: 891, dtype: object
```

```
# Filling missing Age values with mean
train['Age'] = train['Age'].fillna(train['Age'].mean())
# Filling missing Embarked values with most common value
train["Embarked"] = train["Embarked"].fillna(train["Embarked"].mode()[0])
# Converting Pclass to string
train["Pclass"] = train["Pclass"].apply(str)
# Getting Dummies for categorical variables
for col in train.dtypes[train.dtypes == "object"].index:
    for_dummy = train.pop(col)
    train = pd.concat([train, pd.get_dummies(for_dummy, prefix=col)], axis=1)
train.head()
```

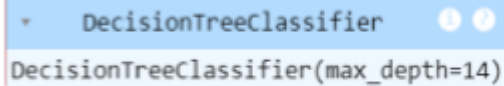
	PassengerId	Survived	Age	SibSp	Parch	Fare	Pclass_1	Pclass_2	Pclass_3	Sex_female	Sex_male	Embarked_C	Embarked_Q	Embarked_S
0	1	0	22.0	1	0	7.2500	False	False	True	False	True	False	False	True
1	2	1	38.0	1	0	71.2833	True	False	False	True	False	True	False	False
2	3	1	26.0	0	0	7.9250	False	False	True	True	False	False	False	True
3	4	1	35.0	1	0	53.1000	True	False	False	True	False	False	False	True
4	5	0	35.0	0	0	8.0500	False	False	True	False	True	False	False	True

```
# Separating labels
labels = train.pop("Survived")

# Splitting the dataset into training and testing sets
x_train, x_test, y_train, y_test = train_test_split(train, labels,
    test_size=0.25)

dt_model = DecisionTreeClassifier(max_depth=14)

dt_model.fit(x_train, y_train)
```



A screenshot of a Jupyter Notebook cell. The cell has a blue header bar with the text "DecisionTreeClassifier" and two circular icons. Below the header, the code "DecisionTreeClassifier(max_depth=14)" is displayed in a monospaced font.

```
y_pred = dt_model.predict(x_test)

accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

_classification_report = classification_report(y_test, y_pred)
print("Classification Report \n", _classification_report)
```

```
Accuracy 0.7399103139013453
Classification Report
      precision    recall  f1-score   support

     0       0.78      0.81      0.80       140
     1       0.66      0.63      0.64        83

   accuracy          0.74       223
  macro avg       0.72      0.72      0.72       223
 weighted avg       0.74      0.74      0.74       223
```

```
boosting_model = AdaBoostClassifier(estimator=DecisionTreeClassifier(),  
n_estimators=50, random_state=42)  
boosting_model.fit(x_train, y_train)  
boosting_predictions = boosting_model.predict(x_test)  
boosting_accuracy = accuracy_score(y_test, boosting_predictions)  
print("Boosting Accuracy:", boosting_accuracy)  
print("Boosting Classification Report \n", classification_report(y_test,  
boosting_predictions))
```

```
Boosting Accuracy: 0.7040358744394619  
Boosting Classification Report  
              precision    recall  f1-score   support  
  
    0           0.77       0.76       0.76       140  
    1           0.60       0.61       0.61        83  
  
 accuracy              0.70       223  
 macro avg           0.68       0.69       0.68       223  
weighted avg           0.71       0.70       0.70       223
```


Practical 18

Aim: Implementation of Gradient Boosting Algorithm

Code:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.metrics import accuracy_score, classification_report
```

```
train = pd.read_csv('titanic.csv') print(train.shape)
```

```
(891, 12)
```

```
# Checking for missing data
NAs = pd.concat([train.isnull().sum()], axis=1, keys=["Train"])
NAs[NAs.sum(axis=1) > 0]# Display columns with missing values
```

Train	
Age	177
Cabin	687
Embarked	2

```
# Removing unnecessary columns
train.pop("Cabin")
train.pop("Name")
train.pop("Ticket")
```

```
0          A/5  21171
1          PC  17599
2    STON/O2. 3101282
3          113803
4          373450
...
886          211536
887          112053
888    W./C.  6607
889          111369
890          370376
Name: Ticket, Length: 891, dtype: object
```

```
# Filling missing Age values with mean
train['Age'] = train['Age'].fillna(train['Age'].mean())

# Filling missing Embarked values with most common value
train["Embarked"] = train["Embarked"].fillna(train["Embarked"].mode()[0])

# Converting Pclass to string
train["Pclass"] = train["Pclass"].apply(str)

# Getting Dummies for categorical variables
for col in train.dtypes[train.dtypes == "object"].index:
    for_dummy = train.pop(col)
train = pd.concat([train, pd.get_dummies(for_dummy, prefix=col)], axis=1)

train.head()
```

PassengerId	Survived	Age	SibSp	Parch	Fare	Pclass_1	Pclass_2	Pclass_3	Sex_female	Sex_male	Embarked_C	Embarked_Q	Embarked_S
0	1	0 22.0	1	0	7.2500	False	False	True	False	True	False	False	True
1	2	1 38.0	1	0	71.2833	True	False	False	True	False	True	False	False
2	3	1 26.0	0	0	7.9250	False	False	True	True	False	False	False	True
3	4	1 35.0	1	0	53.1000	True	False	False	True	False	False	False	True
4	5	0 35.0	0	0	8.0500	False	False	True	False	True	False	False	True

```
# Separating labels
labels = train.pop("Survived")

# Splitting the dataset into training and testing sets
x_train, x_test, y_train, y_test = train_test_split(train, labels,
    test_size=0.25)

# gradient boosting
gb_classifier = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1,
    max_depth=3, subsample=0.8, random_state=42)
gb_classifier.fit(x_train, y_train)
y_pred = gb_classifier.predict(x_test)
gb_accuracy = accuracy_score(y_test, y_pred)
print("Gradient Boosting Accuracy:", gb_accuracy)
print("Gradient Boosting Report \n", classification_report(y_test, y_pred))
```

```
Gradient Boosting Accuracy: 0.8251121076233184
Gradient Boosting Report
      precision    recall  f1-score   support

     0       0.87      0.84      0.86       139
     1       0.75      0.80      0.77        84

 accuracy          0.83       223
  macro avg       0.81      0.82      0.82       223
 weighted avg       0.83      0.83      0.83       223
```